

The cqueues User Guide
for composing
Socket, Signal, Thread, & File Change
Messaging
on
Linux, OS X, Solaris,
FreeBSD, NetBSD, & OpenBSD
with



William Ahern

April 30, 2014

Contents

Contents	i
1 Dependencies	1
1.1 Operating Systems	1
1.1.1 \rightarrow Microsoft Windows	1
1.2 Libraries	1
1.2.1 LuaJIT, Lua 5.2, Lua 5.3	1
1.2.2 OpenSSL	1
1.2.3 pthreads	2
1.3 Compilers	2
1.4 GNU Make	2
2 Installation	3
2.1 Building	3
2.1.1 Targets	3
2.2 Installing	3
2.2.1 Targets	4
3 Usage	5
3.1 Conventions	5
3.1.1 Polling	5
object:pollfd	5
object:events	5
object:timeout	5
3.1.2 \rightarrow Globals	6
3.1.3 Errors	6
EAGAIN	6
ETIMEDOUT	7
EPIPE	7
EBADF	7
3.2 Modules	7
3.2.1 cqueues	7
cqueues.VENDOR	7
cqueues.VERSION	7
cqueues.COMMIT	7
cqueues.interpose	7
cqueues.monotime	7
cqueues.cancel	8

cqueues.poll	8	cqueues:loop	9
cqueues.sleep	8	cqueues:errors	9
cqueues.running	8	cqueues:empty	9
cqueues.new	8	cqueues:count	9
cqueues:attach	8	cqueues:cancel	9
cqueues:wrap	8	cqueues:pause	9
cqueues:step	8		
3.2.2 cqueues.socket			9
socket[]	9	socket:setmaxline	13
socket.interpose	10	socket:settimeout	13
socket.connect	10	socket:onerror	13
socket.connect	10	socket:read	13
socket.listen	10	socket:write	14
socket.listen	10	socket:flush	14
socket.pair	11	socket:fill	14
socket.setvbuf	11	socket:unget	14
socket.setmode	11	socket:pending	14
socket.setbufsiz	11	socket:uncork	14
socket.setmaxline	11	socket:recv	14
socket.settimeout	11	socket:send	15
socket.onerror	11	socket:recvfd	15
socket:connect	11	socket:sendfd	15
socket:listen	11	socket:shutdown	15
socket:accept	11	socket:eof	15
socket:clients	12	socket:peername	15
socket:starttls	12	socket:peerid	16
socket:checktls	12	socket:peerpid	16
socket:setvbuf	12	socket:localname	16
socket:setmode	12	socket:stat	16
socket:setbufsiz	12	socket:close	16
3.2.3 cqueues.errno			16
errno[]	16	errno.strerror	16
3.2.4 cqueues.signal			16
signal[]	16	signal.unblock	17
signal.strsignal	17	signal.raise	17
signal.ignore	17	signal.interpose	17
signal.default	17	signal.listen	17
signal.discard	17	signal:wait	18
signal.block	17	signal:settimeout	18
3.2.5 cqueues.thread			18

thread.self	18	thread:join	18
thread.start	18		
3.2.6 cqueues.notify	19		
notify[]	19	notify:add	19
notify.flags	19	notify:get	19
notify.opendir	19	notify:changes	19
3.2.7 cqueues.dns.record	19		
record.type[]	20	record:name	20
record.class[]	20	record:type	20
record.sshfp[]	20	record:class	20
record:section	20	record:ttd	20
3.2.8 cqueues.dns.packet	20		
packet.section[]	21	packet:qid	21
packet.opcode[]	21	packet:flags	21
packet.rcode[]	21	packet:count	22
packet.interpose	21	packet:grep	22
packet.new	21		
3.2.9 cqueues.dns.config	22		
config[]	22	config:loadfile	23
config.interpose	22	config:loadpath	23
config.new	22	config:get	23
config.stub	23	config:set	23
config.root	23		
3.2.10 cqueues.dns.hosts	24		
hosts.interpose	24	hosts:loadfile	24
hosts.new	24	hosts:loadpath	24
hosts.stub	24	hosts:insert	24
hosts.root	24		
3.2.11 cqueues.dns.hints	24		
hints.interpose	24	hints.root	25
hints.new	25	hints:insert	25
hints.stub	25		
3.2.12 cqueues.dns.resolver	25		

resolver.interpose	25	resolver:query	26
resolver.new	25	resolver:stat	26
resolver.stub	25	resolver:close	26
resolver.root	25		
3.2.13 cqueues.condition	26		
condition.type	26	condition:wait	27
condition.interpose	27	condition:signal	27
condition.new	27		
4 Examples			28
4.1 HTTP SSL Request			28
4.2 Multiplexing Echo Server			29
4.3 Thread Messaging			30

1 Dependencies

1.1 Operating Systems

`cqueues` heavily relies on a modern POSIX environment. But the fundamental premise is to build on the new but non-standard polling facilities provided by contemporary Unix environments. Specifically, BSD `kqueue`, Linux `epoll`, and Solaris Event Ports.

`cqueues` should work on recent versions of Linux, OS X, Solaris, NetBSD, FreeBSD, OpenBSD, and derivatives. The only other possible candidate is AIX, if and when support for AIX's `pollset` interface is added to the embedded “`kpoll`” library.

1.1.1 \neg Microsoft Windows

Microsoft Windows support is basically out of the question¹, for far too many reasons to put here. Aside from the more technical reasons, Windows I/O and networking programming interfaces have a fundamentally different character than on Unix. Unix historically relies on readiness polling, while Windows uses event completion callbacks. There are strengths and weaknesses to each approach. Trying to paper over the chasm between the two approaches invariably results in a framework with the strengths of neither and the weaknesses of both. The purpose of `cqueues` is to leverage the strengths of polling as well as address the weaknesses.

1.2 Libraries

1.2.1 LuaJIT, Lua 5.2, Lua 5.3

`cqueues` principally targets Lua 5.2 and above. It's not fully portable to Lua 5.1 because `cqueues` relies on ephemeron tables to prevent coroutine/controller reference cycles, and because Lua 5.1 does not support yielding from metamethods and iterators. LuaJIT removes the latter of these handicaps, and so `cqueues` targets LuaJIT secondarily. In lieu of ephemeron tables, application code must be sure not to hold a reference to a parent controller in an upvalue of the coroutine. Instead, use `cqueues.running`.

1.2.2 OpenSSL

The `cqueues socket` module provides seamless SSL/TLS support using OpenSSL.

¹I have been toying with the idea of using an `fd_set` in-place of a pollable descriptor on Windows, and taking the union of all `fd_sets` when polling.

Comprehensive bindings for certificate and key management are provided in the [companion openssl module, luaossl](#).

1.2.3 pthreads

`cqueues` provides an optional threading module, using POSIX threads.² Internally it consistently uses thread-safe routines when built with either the `_REENTRANT` or `_THREAD_SAFE` feature macros, such as `pthread_sigmask` instead of `sigprocmask`. Thread support is enabled by default.

Linking Note that on some systems, such as NetBSD and FreeBSD, the loading application must be linked against pthreads (using `-lpthread` or `-pthread`). It is not enough for the `cqueues` module to pull in the dependency at load time. In particular, if using the stock Lua interpreter, it must have been linked against pthreads at build time. Add the appropriate linker flag to MYLIBS in `lua-5.2.x/src/Makefile`.

OpenBSD OpenBSD 5.1 threading is completely *fubar*, especially with regard to signals, because of OpenBSD's transition to kernel threading. If using OpenBSD, be sure to compile *without* the thread-safe macros predefined, especially if using `cqueues.signal`.

1.3 Compilers

The source code is mostly ISO C99 compliant, and even more so with regards to ISO C11. But regardless of standards conformance, it aims to build cleanly with the native compiler for each targeted platform. It currently builds with recent versions of GCC, clang, and SunPro.

Patches are welcome to silence compiler diagnostics.

1.4 GNU Make

The Makefile requires GNU Make, usually installed as `gmake` on platforms other than Linux or OS X. The actual Makefile proxies to `GNUmakefile`. As long as `gmake` is installed on non-GNU systems you can invoke your system's `make`.

²Building without threading enabled is not well tested.

2 Installation

All the C modules are built into a single core C library. The core routines are then wrapped and extended through Lua modules. Because there several extant versions of Lua often used in parallel on the same system, there are individual targets to build and install for each supported Lua version. The targets `all` and `install` will attempt to build and install both Lua 5.1 and 5.2 modules.

Note that building and installation can be accomplished in a single step by simply invoking one of the install targets with all the necessary variables defined.

2.1 Building

There is no separate `./configure` step. System introspection occurs during compile-time. However, the “`configure`” make target can be used to cache the build environment so one needn’t continually use a long command-line invocation.

All the common GNU-style compiler variables are supported, including `CC`, `CPPFLAGS`, `CFLAGS`, `LDFLAGS`, and `SOFLAGS`. Note that you can specify the path to Lua 5.1, Lua 5.2, and Lua 5.3 include headers at the same time in `CPPFLAGS`; the build system will work things out to ensure the correct headers are loaded when compiling each version of the module.

2.1.1 Targets

`all`

Build modules for Lua 5.1 and 5.2.

`all5.1`

Build Lua 5.1 module.

`all5.2`

Build Lua 5.2 module.

`all5.3`

Build Lua 5.3 module.

2.2 Installing

All the common GNU-style installation path variables are supported, including `prefix`, `bindir`, `libdir`, `datadir`, `includedir`, and `DESTDIR`. These additional path variables are also allowed:

lua51path

Install path for Lua 5.1 modules, e.g. `$(prefix)/share/lua/5.1`

lua51cpath

Install path for Lua 5.1 C modules, e.g. `$(prefix)/lib/lua/5.1`

lua52path

Install path for Lua 5.2 modules, e.g. `$(prefix)/share/lua/5.2`

lua52cpath

Install path for Lua 5.2 C modules, e.g. `$(prefix)/lib/lua/5.2`

lua53path

Install path for Lua 5.3 modules, e.g. `$(prefix)/share/lua/5.3`

lua53cpath

Install path for Lua 5.3 C modules, e.g. `$(prefix)/lib/lua/5.3`

2.2.1 Targets

install

Install modules for Lua 5.1 and 5.2.

install5.1

Install Lua 5.1 module.

install5.2

Install Lua 5.2 module.

install5.3

Install Lua 5.3 module.

3 Usage

3.1 Conventions

3.1.1 Polling

`cqueues` works through a simple protocol. When a coroutine yields to its parent `cqueues` controller, it can pass one or more objects. These objects are introspected for three methods: `:pollfd`, `:events`, and `:timeout`. These methods generate the parameters for installing descriptor and timeout events. When one of these events fires, `cqueues` will resume the coroutine, passing the relevant objects which were interested in the triggered event. It's analogous to calling `poll`, and in fact the routine `cqueues.poll` is provided as a wrapper for `coroutine.yield`.¹

`:pollfd()`

The `:pollfd` method should return a descriptor integer or `nil`. This descriptor must remain in existence until the owner object is garbage collected, `cqueues.cancel` is used, the coroutine executes one additional yield/resume cycle (so the old descriptor is expired from the descriptor queue), or until after the coroutine exits. If the descriptor is closed prematurely, the kernel will remove it from the internal descriptor queue, bringing it out of sync with the controller, and probably causing `cqueues:step` to return `EBADF` or `ENOENT` errors.

`:events()`

The `:events` method should return a string or `nil`. `cqueues` searches the string for the flags 'r' and 'w', which describe the events to associate with the descriptor—respectively, `POLLIN` and `POLLOUT`.

`:timeout()`

The `:timeout` should return a number or `nil`. This schedules an independent timeout event. To effect a simple one second timeout, you can do

```
cqueues.poll({ timeout = function() return 1.0 end })
```

which is equivalent to

```
coroutine.yield({ timeout = function() return 1.0 end })
```

¹In the future this wrapper may be able to detect if the current coroutine was resumed by a controller, and if not chain yield calls until a controller is reached.

but not quite as simple as the shortcut

```
cqueues . poll (1.0)
```

Instantiated `cqueues` objects implement all three methods.² In particular, this means that you can stack `cqueues`, or poll on a `cqueues` object using some other event loop library. Each `cqueues` object is entirely self-contained, without any global state.

3.1.2 \neg Globals

Like the core controller module, other `cqueues` modules adhere to a *no global side effects* discipline. In particular, this means

- no global process variables;
- no signal handling gimmicks—like the pipe trick—which could conflict with other components of your application³;
- consistent use of thread-safe function variants; and
- consistent use of `O_CLOEXEC` and similar flags to eliminate or reduce `fork` + `exec` races in threaded applications.

3.1.3 Errors

The usual behavior is for errors to be returned directly. But see `socket.onerror`. If a routine is specified to return an object or string, `nil` is returned; if a boolean, `false` is returned. In both cases, these are usually followed by a numeric error code. Thus, if a routine is specified to return two values on success, then on error three values are returned, the first two `nil` or `false`, and the third an error code.

`cqueues` is a relatively low-level component library. In almost all cases errors will be system errors, returned as numeric error codes for easy and efficient comparison. For example, attempting to create a UNIX domain socket with `socket.listen` in a directory without sufficient permissions might return `‘nil, EACCES’`.

EAGAIN

`cqueues` modules are implemented in both C and Lua. The C routines never yield, and always return recoverable errors directly. Most C routines are wrapped—and methods interposed—with Lua functions. These Lua functions usually poll when `EAGAIN` is encountered and retry the C routine on resumption. Few methods will return `EAGAIN` directly.

²:`pollfd` returns the internal `kqueue`, `epoll`, or Ports descriptor; `:events` returns “r”; and `:timeout` returns the time to the next internal timeout event.

³The `cqueues.thread` module ensures threads are started with a filled signal mask.

ETIMEDOUT

This error value is usually seen when a timeout is specified by the caller of a logically synchronous method. The method will normally yield and poll if the operation cannot be completed immediately, but if the timeout expires then it will return a failure with `ETIMEDOUT`.

EPIPE

In Unix `EPIPE` is only encountered when attempting to write to a closed pipe or socket. In `cqueues` `EPIPE` is used to signal both EOF and a closed output stream.⁴ The low-level I/O method `socket:recv`, for example, returns `EPIPE` on EOF. In other cases, as with `socket:read`, EOF is not an error condition.

EBADF

This error commonly occurs in asynchronous applications, which are especially prone to bugs related to their complex state management. With Lua code using the `cqueues` APIs, `EBADF` should never be encountered. When it does occur, it's a sure sign of a bug somewhere in the parent application or an extension module and—hopefully—not `cqueues`.

3.2 Modules

3.2.1 `cqueues`

`cqueues.VENDOR`

String describing the vendor, e.g. `william@25thandClement.com`. If you fork this project please change this string so I don't receive unwarranted scorn or praise.

`cqueues.VERSION`

Number describing the running version, formatted as `YYYYMMDD`. Official releases are tagged in the git repo as `rel-YYYYMMDD`.

`cqueues.COMMIT`

Git commit hash string of HEAD.

`cqueues.interpose(name, function)`

Add or interpose a `cqueues` controller class method. Returns the previous method, if any.

`cqueues.monotime()`

Return the system's monotonic clock time, usually `clock_gettime(CLOCK_MONOTONIC)`.

⁴In some situations, such as with SSL/TLS, a read attempt might require a write, anyhow. Expanding the scope of `EPIPE` simplifies the logic required to handle various I/O failures.

`cqueues.cancel(fd)`

Cancels the specified descriptor for all controllers. This ensures safe early closure of descriptors. However, the complexity is approximately $M \cdot 2 \log N$, where M is the number of controllers, and N the number of descriptors per controller (presuming equal distribution). For most purposes this is entirely inconsequential. By contrast, however, implicit cancellation through GC or yield/resume cycling is $O(1)$.

Any coroutine polling on the canceled descriptor is placed on its controller's pending queue.

`cqueues.poll(...)`

Takes a series of objects obeying the polling protocol and yields control to the parent `cqueues` controller. On an event resumes the coroutine, passing the objects which triggered resumption. A number value is interpreted as a timeout.

`cqueues.sleep(number)`

Yields to the parent `cqueues` controller and schedules a wakeup for 'number' seconds in the future.

`cqueues.running()`

Returns two values: the immediate controller currently executing, if any, or nil; and a boolean—true if the caller's coroutine is the same coroutine resumed by the controller.

`cqueues.new()`

Create a new `cqueues` object.

`cqueue:attach(coroutine)`

Attach and manage the specified coroutine.

`cqueue:wrap(function)`

Execute function inside a new coroutine managed by the controller.

`cqueue:step([timeout])`

Step once through the event queue. Unless the timeout is explicitly specified as 0, or unless the current thread of execution is a `cqueues` managed coroutine, *it suspends the process indefinitely or for the specified timeout* until a descriptor event or timeout fires.

Returns true or false. If false—i.e. a coroutine exited abnormally—then a second return value holds the error message. `:step` can be called again after errors.

If embedding `cqueues` within an existing application, the top-level `:step` invocation should always specify a 0 timeout. A controller is a pollable object, and the descriptor returned by the `:pollfd` method can be used with third-party event libraries, whether written in Lua, C, or some other language. Don't forget to also schedule a timeout using the value from `:timeout`.

`cqueue:loop([timeout])`

Invoke `cqueues:step` in a loop, exiting on error, timeout, or if the event queue is empty. Returns true if no error occurred, or false and an error string from `cqueues:step`.

`cqueue:errors([timeout])`

Returns an iterator function over errors returned from `cqueues:loop`. If `cqueues:loop` returns successfully because of an empty event queue, or if the timeout expires, returns nothing, which terminates any for-loop. ‘timeout’ is cumulative over the entire iteration, not simply passed as-is to each invocation of `cqueues:loop`.

`cqueue:empty()`

Returns true if there are no more descriptor or timeout events queued, false otherwise.

`cqueue:count()`

Returns a count of managed coroutines.

`cqueue:cancel(fd)`

Cancel the specified descriptor for that controller. See `cqueues.cancel`.

`cqueue:pause(signal [, signal ...])`

A wrapper around `pselect` which *suspends execution of the process* until the controller polls ready or a signal is delivered. This interface is provided as a very basic least common denominator for simple slave process controller loops and similar scenarios, where immediate response to signal delivery is required on platforms like Solaris without a proper signal polling primitive. (`signal.listen` on Solaris merely periodically queries the pending set.)

Much better alternatives are possible for Solaris, but require global process state or an LWP thread helper.

3.2.2 `cqueues.socket`

The socket bindings provide built-in DNS, SSL/TLS, buffering, and line translation. DNS happens transparently, and SSL/TLS can be initiated with the `socket:starttls` method.

The default I/O mode is “tl”—text translation and line buffering. This makes sockets work intuitively with the most common protocols on the Internet, like SMTP and HTTP, which require CRLF and use line delimited framing.

`socket[]`

A table mapping socket related system identifier names to number codes, including `AF_UNSPEC`, `AF_INET`, `AF_INET6`, `AF_UNIX`, `SOCK_STREAM`, and `SOCK_DGRAM`.

`socket.interpose(name, function)`

Add or interpose a socket class method. Returns the previous method, if any.

`socket.connect(host, port [, family] [, type])`

Return a new socket immediately ready for reading or writing. DNS lookup and TCP connection handling are handled transparently.

`socket.connect{ ... }`

Like `socket.connect` with list arguments, but takes a table of named arguments:

field	type:default	description
<code>.host</code>	string:nil	IP address or host domain name
<code>.port</code>	string:nil	host port
<code>.path</code>	string:nil	UNIX domain socket path
<code>.family</code>	number	protocol family— <code>AF_INET</code> (default), <code>AF_INET6</code> , <code>AF_UNIX</code> (default if <code>.path</code> specified)
<code>.type</code>	number	protocol type— <code>SOCK_STREAM</code> (default) or <code>SOCK_DGRAM</code>
<code>.mode</code>	string:nil	<code>fchmod</code> or <code>chmod</code> socket after creating UNIX domain socket
<code>.mask</code>	string:nil	set and restore umask when binding UNIX domain sockets
<code>.unlink</code>	boolean:false	unlink socket path before binding
<code>.reuseaddr</code>	boolean:true	<code>SO_REUSEADDR</code> socket option
<code>.nodelay</code>	boolean:false	<code>TCP_NODELAY</code> IP option
<code>.nopush</code>	boolean:false	<code>TCP_NOPUSH</code> , <code>TCP_CORK</code> , or equivalent IP option
<code>.nonblock</code>	boolean:true	<code>O_NONBLOCK</code> descriptor flag
<code>.cloexec</code>	boolean:true	<code>O_CLOEXEC</code> descriptor flag
<code>.nosigpipe</code>	boolean:true	<code>O_NOSIGPIPE</code> , <code>SO_NOSIGPIPE</code> , <code>MSG_NOSIGNAL</code> , or equivalent descriptor flag
<code>.verify</code>	boolean:false	require SSL certificate verification
<code>.time</code>	boolean:true	track elapsed time for statistics

`socket.listen(host, port)`

Return a new socket immediately ready for accepting connections.

`socket.listen{ ... }`

Like `socket.listen` with list arguments, but takes a table of named arguments. See also `socket.connect{}`.

`socket.pair([type])`

Returns two bound sockets. Type should be the system type number, e.g. `socket.SOCK_STREAM` or `socket.SOCK_DGRAM`.

`socket.setvbuf(mode [, size])`

Set the default output buffering mode for all new sockets. See `socket:setvbuf`.

`socket.setmode([input] [, output])`

Set the default I/O modes for all new sockets. See `socket:setmode`.

`socket.setbufsiz([input] [, output])`

Set the default I/O buffer sizes for all new sockets. See `socket:setbufsiz`.

`socket.setmaxline([input] [, output])`

Set the default I/O line-buffering limits for all new sockets. See `socket:setmaxline`.

`socket.settimeout([timeout])`

Set the default timeout for all new sockets. See `socket:setmaxline`.

`socket.onerror([function])`

Set the default error handler for all new sockets. See `socket:onerror`.

`socket:connect([timeout])`

Wait for connection establishment to succeed. You do not need to wait before proceeding to perform read or write calls, but waiting may ease diagnosing connection problems in your code and allows you to separate connect phase from I/O phase timeouts.

`socket:listen([timeout])`

Wait for socket binding to succeed. You do not need to wait before proceeding to call `:accept`, but waiting may ease diagnosing binding problems in your code and allows you to separate listen phase from accept phase timeouts.

Socket binding may not occur immediately if you provided a host address that required DNS resolution over the network. This is uncommon for listening sockets but supported nonetheless; the symmetry simplifies internal code. Also, socket object instantiation with `socket.listen` and `socket.connect` only return errors regarding user data object construction; address lookup and binding errors are detected later, when initiated by subsequent method calls.

`socket:accept([timeout])`

Wait for and return an incoming client socket on a listening object.

`socket:clients([timeout])`

Iterator over `socket:accept:` for `con` in `srv:clients()` do ... end.

`socket:starttls([context][, timeout])`

Place socket into TLS mode, optionally using the `openssl.ssl.context` object as the configuration prototype, and wait for the handshake to complete.⁵ Returns true on success, false and an error code on failure.

`socket:checktls()`

If in TLS mode, returns an `openssl.ssl` object, otherwise nil. If the `openssl` module cannot be loaded, returns nil and an error string.

`socket:setvbuf(mode [, size])`

Same as Lua `file:setvbuf`. Analogous to “n”, “l”, and “f” mode flags. Returns the previous output mode and output buffer size.

`socket:setmode([input] [, output])`

Sets the the input and output buffering and translation modes. Either mode can be nil or none, in which case the mode is left unchanged.

A mode is specified as a string containing one or more of the following flags

flag	description
t	text mode; input or output undergoes LF/CRLF translation
b	binary mode; no LF/CRLF translation
n	no output buffering
l	line buffered output
f	fully buffered output

Returns the previous input and output modes as fixed-sized strings. At present the first character is one of “t” or “b”, and the second character one of “n”, “l”, “f”, or “-” (for in the input mode).

`socket:setbufsiz([input] [, output])`

Sets the input and output buffer size. Either size can be nil or none, in which case the size is left unchanged.

These are not hard limits for `SOCK_STREAM` sockets. The input buffer argument simply sets a minimum for input buffering, to reduce syscalls. The output buffer argument is the same as

⁵Prior to 2014-04-30, if no timeout was specified then the routine returned immediately.

provided to `:setvbuf`, and effectively changes when flushing occurs for full- or line-buffered output modes.

For `SOCK_DGRAM` sockets, the input buffer sets a hard limit on the size of datagram messages. Any message over this size will be truncated, unless a previous block- or line-buffered read operation forced the buffer to be reallocated to a larger size.

Returns the previous input and output buffer sizes, or throws an error if the buffers could not be reallocated.

`socket:setmaxline([input] [, output])`

Sets the maximum input and output length for line-buffered operations. Either size can be nil or none, in which case the size is left unchanged.

These are hard limits. For line-buffered input operations, if a `\n` character is not found within this limit then the data is processed as-if EOF was reached at this boundary. For line-buffered output, a chunk is always flushed at this boundary.

Returns the previous input and output sizes.

`socket:settimeout([timeout])`

Sets the default timeout period for I/O. If nil or none, then clears any default timeout. If a timeout is cleared, any operation which polls will wait indefinitely until completion or an error occurs.

Sockets are instantiated without a default timeout.

`socket:onerror([function])`

Set the error handler. The error handler is passed the tuple socket-object, method-string, error-number, and is expected to either throw an error or return an error code—to be returned to the caller as part of the documented return interface.

The default error handler returns `EPIPE` and `ETIMEDOUT` directly, and throws everything else. `EAGAIN` is handled internally for logically synchronous calls.

Returns the previous error handler, if any.

`socket:read(...)`

Similar to Lua's `file:read`, with additional formats.

format	description
*n	unsupported
*a	read until EOF
*l	read the next line, trimming the EOL marker
*L	read the next line, keeping the EOL marker
number	read ‘number’ bytes or until EOF
*h	read and unfold MIME compliant header
*H	read MIME compliant header, keeping EOL markers
-number	read 1 to ‘number’ bytes, immediately returning if possible

Note that for `SOCK_DGRAM` sockets, each message is treated as-if EOF was reached. The slurp operation returns a single datagram, and line-buffered operations will return the remaining text in a message even without a terminating `\n`. Datagrams will be truncated if the message is larger than the input buffer size.

`socket:write(...)`

Same as `Lua file:write`.

`socket:flush([mode] [, timeout])`

Flushes output buffer. Mode is one of the “nlf” flags described in `socket.connect`. A nil mode implies “n”, i.e. no buffering and effecting a full flush. An empty string mode resolves to the configured output buffering mode.

`socket:fill(size[, timeout])`

Fills the input buffer with ‘size’ bytes. Returns true on success, false and an error code on failure.

`socket:unget(string)`

Writes ‘string’ to the head of the socket input buffer.

`socket:pending()`

Returns two numbers—the counts of buffered bytes in the input and output streams. This does not include the bytes in the kernel’s buffer.

`socket:uncork()`

Disables `TCP_NOPUSH`, `TCP_CORK`, or equivalent socket option.

`socket:recv(format [, mode])`

Similar to `socket:read`, except takes only a single format and returns immediately without polling. On success returns the string or number. On failure returns nil and a numeric error code—usually `EAGAIN` or `EPIPE`. Does not use error handler.

‘mode’ is as described in `socket.connect`, and defaults to the configured input mode.

`socket:send(string, i, j [, mode])`

Write out the slice `‘string’[i,j]`. Similar to passing `string:sub(i, j)`, but without instantiating a new string object. Immediately returns two values: count of bytes written (0 to `j-i+1`), and numerical error code, if any (usually `EAGAIN` or `EPIPE`).

`socket:recvfd([prebufsiz] [, timeout])`

Receive an ancillary socket message with accompanying descriptor. ‘prebufsiz’ specifies the maximum message size to expect.

This routine bypasses I/O buffering.

Returns message-string, socket-object on success; nil, nil, error-integer on failure. On success socket-object may still be nil. Message truncation is treated as an error condition.

`socket:sendfd(msg, socket[, timeout])`

Send an ancillary socket message with accompanying descriptor. ‘msg’ should be a non-zero-length string, which some platforms require. ‘socket’ should be a Lua file handle, `cqueues` socket, integer descriptor, or nil.

This routine bypasses I/O buffering.

Returns true on success; false and an error code on failure.

`socket:shutdown(how)`

Simple binding to `shutdown(2)`. ‘how’ is a string containing one or both of the flags “r” or “w”.

flag	description
r	analagous to <code>shutdown(SHUT_RD)</code>
w	analagous to <code>shutdown(SHUT_WR)</code>

`socket:eof()`

Returns two boolean values representing whether EOF has been received on the input channel, and whether the output channel has signaled closure (e.g. `EPIPE`).

`socket:peername()`

Returns one, two, or three values. On success, returns three values for `AF_INET` and `AF_INET6` sockets—the address family number, IP address string, and IP port. For `AF_UNIX` sockets, returns the address family and file path. If the socket is not yet connected, returns the address family `AF_UNSPEC`, usually numeric 0.

On failure returns nil and a numeric error code.

`socket:peereid()`

Queries the effective UID and effective GID of an AF_UNIX, SOCK_STREAM peer as cached by the kernel when the stream initially connected.

Returns two numbers representing the UID and GID, respectively, on success, otherwise nil and a numeric error code.

`socket:peerpid()`

Queries the PID of a AF_UNIX, SOCK_STREAM peer as cached by the kernel when the stream initially connected. This capability is unsupported on OS X and FreeBSD; they only provide `getpeereid`, which cannot provide the PID.

Returns a number representing the PID on success, otherwise nil and a numeric error code.

`socket:localname()`

Identical to `socket:peername`, but returns the local address of the socket.

`socket:stat()`

Returns a table containing two subtables, ‘sent’ and ‘rcvd’, which each have three fields—.count for the number of bytes sent or received, a boolean .eof signaling whether input or output has been shutdown, and .time logging the last send or receive operation.

`socket:close()`

Explicitly and immediately close all internal descriptors. This routine ensures all descriptors are properly cancelled.

3.2.3 `cqueues.errno`

`errno[]`

A table mapping all system error string macros to numerical error codes, and all numerical error codes to system error string macros. Thus, `errno.EAGAIN` evaluates to a numeric error code, and `errno[errno.EAGAIN]` evaluates to the string “EAGAIN”.

`errno.strerror(code)`

Returns string returned by `strerror(3)`.

3.2.4 `cqueues.signal`

`signal[]`

A table mapping signal string macros to numerical signal codes. In all likelihood, `signal.SIGKILL` evaluates to the number 9.

`signal.strsignal(code)`

Returns string returned by `strsignal(3)`.

`signal.ignore(signal [, signal ...])`

Set the signal handler to `SIG_IGN` for the specified signals.

`signal.default(signal [, signal ...])`

Set the signal handler to `SIG_DFL` for the specified signals.

`signal.discard(signal [, signal ...])`

Set the signal handler to a builtin “noop” handler for the specified signals. Use this if you want signals to interrupt syscalls.

`signal.block(signal [, signal ...])`

Block the specified signals.

`signal.unblock(signal [, signal ...])`

Unblock the specified signals.

`signal.raise(signal [, signal ...])`

`raise(3)` the specified signals.

`signal.interpose(name, function)`

Add or interpose a signal listener class method. Returns the previous method, if any.

`signal.listen(signal [, signal ...])`

Returns a signal listener object for the specified signals. Semantics differ between platforms:

kqueue BSD `kqueue` provides the most intuitive behavior. All listeners will detect a signal sent to the process irrespective of whether the signal is ignored, blocked, or delivered. However, `EVFILT_SIGNAL` is edge-triggered, which means no notification of delivery of a pending signal upon being unblocked.

signalfd Linux `signalfd` will not detect ignored or delivered signals, and only one `signalfd` object will poll ready per signal.

sigtimedwait Solaris provides no signal polling kernel primitive. Instead, the pending set is periodically queried using `sigtimedwait`. See `signal:settimeout`. Like Linux, only one listener can notify per interrupt.

To be portable the application must block the relevant signals. See `signal.block`. Otherwise, neither Linux nor Solaris will be able to detect the interrupt. Any signal should be assigned to one listener only, although any listener may query multiple signals.

Alternatively, applications may start a dedicated thread to field incoming signals, and send notifications over a socket. In the future this may be provided as an optional listener implementation.

See also `cqueue:pause` for another, if crude, alternative.

`signal:wait([timeout])`

Polls for the signal set passed to the constructor. Returns the signal number, or nil on timeout.

`signal:settimeout(timeout)`

Set the polling interval for implementations such as Solaris which lack a signal polling kernel primitive. On such systems `signal:wait` merely queries the pending set every ‘timeout’ seconds.

3.2.5 `cqueues.thread`

`thread.self()`

Returns the LWP thread object for the running Lua instances. Threads not started via `thread.start` return nil.

`thread.start(function [, string [, string ...]])`

Generates a socket pair, starts a POSIX LWP thread, initializes a new Lua VM instance, preloads the `cqueues` library, and loads and executes the specified function from the new LWP thread and Lua instance. The function receives as the first parameter one end of the socket pair—instantiated as a `cqueues.socket` object—followed by the string parameters passed to `thread.start`.

The new LWP thread starts with all signals blocked.

Returns a thread object and a socket object—the other end of the socket pair. The thread object is pollable, and readiness signals that the LWP thread has exited, or is imminently about to exit.

On error returns two nils and an error code.

`thread.join([timeout])`

Wait for the thread to terminate. Calling the equivalent of `thread.self():join()` is disallowed.

Returns a boolean and error value. If false, error value is an error code describing a local error, usually `EAGAIN` or `ETIMEDOUT`. If true, error value is 1) an error code describing a system error which the thread encountered, 2) an error message string returned by the new Lua instance, or 3) nil if completed successfully.

3.2.6 `cqueues.notify`

`notify[]`

A table mapping bitwise flags to names, and vice-versa.

name	description
CREATE	file creation event
ATTRIB	metadata change event
MODIFY	modification to file contents or directory entries
REVOKE	permission revoked
DELETE	file deletion event
ALL	bitwise-or of CREATE, DELETE, ATTRIB, MODIFY, and REVOKE

`notify.flags(bitset[, bitset ...])`

Returns an iterator over the flags in the specified bitwise change sets. Thus, `notify.flags(bit32.xor(notify.CREATE, notify.DELETE), notify.MODIFY)` returns an iterator returning all three flags.

`notify.opendir(path[, changes])`

Returns a notification object associated with the specified directory. Directory change events are limited to the set, ‘changes’, or to `notify.ALL` if nil.

`notify.add(name[, changes])`

Track the specified file name within the notification directory. ‘changes’ defaults to `notify.ALL` if nil.

`notify.get([timeout])`

Returns a bitwise change set and a filename on success.

`notify.changes([timeout])`

Returns an iterator over the `notify.get` method.

3.2.7 `cqueues.dns.record`

DNS resource record objects are implemented within `cqueues.dns.record`. The global tables and shared methods are documented below. The type-specific accessory methods are quite numerous. Until documented please confer with `cqueues/src/dns.c`. Also, the accessory method names are usually equivalent to the structure member names in `cqueues/src/lib/dns.h`, which in return usually reflect the member names in the relevant RFC.

The `__toString` metamethod returns a representation of the record data only, excluding the name, type, ttl, etc. For an A record, it’s equivalent to `string.format(“%s”, rr:addr())`. For MX—which has multiple members—it’s `string.format(“%d %s”, rr:preference(), rr:host())`.

`record.type[]`

A table mapping DNS record type string identifiers to number values, and vice-versa. So, `record.type.A` evaluates to 1, the IANA numeric record type. String identifiers are only provided for record types which are directly parseable and composable by the library. Currently supported types include A, NS, CNAME, SOA, PTR, MX, TXT, AAAA, SRV, OPT, SSHFP, and SPF. Other record types can be instantiated, but the numeric type must be used and the only methods available operate on the raw rdata.

`record.class[]`

A table mapping DNS record class string identifiers to number values, and vice-versa. At present the only class included is IN.

`record.sshfp[]`

A table mapping DNS SSHFP record string identifiers to the number values—RSA, DSA, and SHA1.

`record:section()`

Returns the section identifier from whence the record came, if derived from a packet. Specifically, QUESTION, ANSWER, AUTHORITY, or ADDITIONAL. See `cqueues.dns.packet.section[]`.

`record:name()`

Returns the uncompressed record domain name as a string.

`record:type()`

Returns the numeric record type. If ‘rr’ holds an AAAA record, then the return value of `rr:type()` will compare equal to `record.type.AAAA`.

`record:class()`

Returns the numeric record class. See `record.class[]`.

`record:ttl()`

Returns the record TTL.

3.2.8 `cqueues.dns.packet`

DNS packets are stored in a simple structure encapsulating the raw packet data. One consequence is that packets are append only. Because a packet is composed of four adjacent sections, when building a packet all the information necessary should be at-hand so that records can be appended in order.

The `__tostring` metaclass composes a string similar to the output of the venerable dig utility.

`packet.section[]`

A table mapping packet section string identifiers to number values, and vice-versa. A packet is composed of only four sections: QUESTION, ANSWER, AUTHORITY, and ADDITIONAL.

`packet.opcode[]`

A table mapping packet opcode string identifiers to number values, and vice-versa. The currently mapped opcodes are QUERY, IQUERY, STATUS, NOTIFY, and UPDATE.

`packet.rcode[]`

A table mapping packet rcode string identifiers to number values, and vice-versa. The currently mapped rcodes are NOERROR, FORMERR, SERVFAIL, NXDOMAIN, NOTIMP, REFUSED, YXDOMAIN, YXRRSET, NXRRSET, NOTAUTH, and NOTZONE.

`packet.interpose`

Add or interpose a packet class method. Returns the previous method, if any.

`packet.new([prepbuftsiz])`

Instantiate a new packet object. ‘prepbuftsiz’ is the maximum space available for appending compressed records. For constructing a packet with a single question, the most space possibly necessary is 260—256 bytes for the name, and 2 bytes each for the type and class (a QUESTION record has no TTL or rdata section).

`packet:qid()`

Returns the 16-bit QID value.

`packet:flags()`

Returns a table of packet header flags.

field	type	description
.qr	integer	specifies whether the packet is a query (0) or response (1)
.opcode	number	specifies the query type
.aa	boolean	signals an authoritative answer
.tc	boolean	signals packet truncation
.rd	boolean	signals “recursion desired”
.ra	boolean	signals “recursion available”
.z	boolean	reserved by RFC 1035 and used by other RFCs
.rcode	integer	specifies the response disposition

`packet:count([sections])`

Returns a count of records in the sections specified by the bitwise parameter ‘sections’. Defaults to `packet.section.ALL`, which is the XOR of all four sections.

`packet:grep{ ... }`

Returns a record iterator over the packet according to all the criteria specified by the optional table parameter.

field	description
<code>.section</code>	select records by bitwise AND with the specified sections
<code>.type</code>	select records of this type (not bitwise)
<code>.class</code>	selects records of this class (not bitwise)
<code>.name</code>	select records with this name

3.2.9 `cqueues.dns.config`

The traditional BSD `/etc/resolv.conf` file is the prototype for this module, although it’s also capable of parsing `/etc/nsswitch.conf`. `cqueues.dns.config` objects are used when instantiating new resolver objects, and provide the general options controlling a resolver.

The `__tostring` metamethod composes a string adhering to `/etc/resolv.conf` syntax, with `/etc/nsswitch.conf` alternatives as comments.

`config[]`

A table mapping flag identifiers to number values.

field	description
<code>TCP_ENABLE</code>	fall back to TCP when truncation detected (default)
<code>TCP_ONLY</code>	only use TCP when querying
<code>TCP_DISABLE</code>	do not fall back to TCP
<code>RESOLV_CONF</code>	specifies BSD <code>/etc/resolv.conf</code> input syntax
<code>NSSWITCH_CONF</code>	specifies Solaris <code>/etc/nsswitch.conf</code> input syntax

`config.interpose(name, function)`

Add or interpose a config class method. Returns the previous method, if any.

`config.new{ ... }`

Returns a new config object, optionally initialized according to the specified table values.

field	type	description
<code>.nameserver</code>	table	list of IP address strings to use for stub resolvers
<code>.search</code>	table	list of domain suffixes to append to query names

<code>.lookup</code>	table	order of lookup methods—“file” and “bind”
<code>.options</code>	table	canonical location for <code>.edns0</code> , <code>.ndots</code> , <code>.timeout</code> , <code>.attempts</code> , <code>.rotate</code> , <code>.recurse</code> , <code>.smart</code> , and <code>.tcp</code> options
<code>..edns0</code>	boolean	enable EDNS0 support
<code>..ndots</code>	number	if query name has fewer labels than this, reverse suffix search order
<code>..timeout</code>	number	timeout between query retries
<code>..attempts</code>	number	maximum number of attempts per nameserver
<code>..rotate</code>	boolean	randomize nameserver selection
<code>..recurse</code>	boolean	query recursively instead of as a simple stub resolver
<code>..smart</code>	boolean	for NS, MX, SRV and similar record queries, resolve the A record if not included as glue in the initial answer
<code>..tcp</code>	number	see <code>TCP_ENABLE</code> , <code>TCP_ONLY</code> , <code>TCP_DISABLE</code> in <code>config[]</code>
<code>.interface</code>	string	IP address to bind to when querying (e.g. <code>[192.168.1.1]:1234</code>)

`config.stub{ ... }`

Returns a config object initialized for a stub resolver by loading the relevant system files; e.g. `/etc/resolv.conf` and `/etc/nsswitch.conf`. Takes optional initialization values like `config.new`.

`config.root{ ... }`

Returns a config object initialized for a recursive resolver. Takes optional initialization values like `config.new`.

`config:loadfile(file[, syntax])`

Parse the Lua file object ‘file’. ‘syntax’ describes the format, which should be `RESOLV_CONF` (default), or `NSSWITCH_CONF`.

`config:loadpath(path[, syntax])`

Like `:loadfile`, but takes a file path.

`config:get()`

Returns the configuration as a Lua table structure. See `config.new` for a description of the values.

`config:set{ ... }`

Apply the defined configuration values. The table should have the same structure as described for `config.new`.

3.2.10 `cqueues.dns.hosts`

The traditional BSD `/etc/hosts` file is the prototype for this module, and provides resolvers the data source for the “file” lookup method.

The `__tostring` metamethod composes a string adhering to `/etc/hosts` syntax.

`hosts.interpose(name, function)`

Add or interpose a hosts class method. Returns the previous method, if any.

`hosts.new()`

Returns a new hosts object.

`hosts.stub()`

Returns a host object initialized for a stub resolver by loading the relevant system files; e.g. `/etc/hosts`.

`hosts.root()`

Returns a hosts object initialized for a recursive resolver.

`hosts:loadfile(file)`

Parse the Lua file object ‘file’ for host entries.

`hosts:loadpath(path)`

Like `:loadfile`, but takes a file path.

`hosts:insert(address, name[, alias])`

Inserts a new hosts entry. ‘address’ should be an IPv4 or IPv6 address string, ‘name’ the domain name, and ‘alias’ a boolean—true if ‘name’ is canonical and a valid response for a reverse address lookup.

3.2.11 `cqueues.dns.hints`

The internal DNS library is implemented as a recursive resolver. No matter whether configured as a stub or recursive resolver, when a query is submitted it consults a “hints” database for the initial name servers to contact. In stub mode these would usually be the local recursive, caching name servers, derived from the `cqueues.dns.config` object; in recursive mode, the root IANA name servers.

The `__tostring` metamethod composes a multi-line string indexing SOA zone names and addresses.

`hints.interpose(name, function)`

Add or interpose a hints class method. Returns the previous method, if any.

`hints.new([resconf])`

Returns a new hints object. ‘resconf’ is an optional `cqueues.dns.config` object which in the future may be used to initialize database behavior. Currently it’s unused, and *does not* pre-load the name server list.

`hints.stub([resconf])`

Returns a hints object initialized for a stub resolver. If provided, the initial hints are taken from the `cqueues.dns.config` object, ‘resconf’. Otherwise, the hints are derived from a temporary “stub” config object internally.

`hints.root([resconf])`

Returns a hints object initialized for a recursive resolver. The root name servers are initialized from an internal database compiled into the module. See `hints.new` for the function of the optional ‘resconf’.

`hints.insert(zone, address|resconf[, priority])`

Inserts a new hints entry. ‘zone’ is the domain name which anchors the SOA (e.g. “.”, or “com.”), and ‘address’ the IPv4 or IPv6 of the nameserver. Alternatively, in lieu of a string address a `cqueues.dns.config` object can be specified, and the addresses taken from the nameserver list property. ‘priority’ is used for ordering nameservers in each zone.

IPv4 and IPv6 addresses can optionally contain a port component, e.g. “[2001:503:ba3e::2:30]:123” or “[198.41.0.4]:53”.

3.2.12 `cqueues.dns.resolver`

`resolver.interpose(name, function)`

Add or interpose a resolver class method. Returns the previous method, if any.

`resolver.new([resconf][,hosts][,hints])`

Returns a new resolver object, configured according to the specified config, hosts, and hints objects. ‘resconf’ can be either an object, or a table suitable for passing to `config.new`. ‘hosts’ and ‘hints’, if nil, are instantiated according to the mode—recursive or stub—of the config object.

`resolver.stub{ ... }`

Returns a stub resolver, optionally initialized to the defined config parameters, which should have a structure suitable for passing to `cqueues.dns.config.new`.

`resolver.root{ ... }`

Returns a recursive resolver, optionally initialized to the defined config parameters, which should have a structure suitable for passing to `cqueues.dns.config.new`.

`resolver:query(name[, type][, class][, timeout])`

Query for the DNS resource record with the specified type and class. ‘name’ is the fully-qualified or prefix domain name string. ‘type’ and ‘class’ are numbers corresponding to the IANA assigned numeric identifier, and default to A (0x01) and IN (0x01), respectively. ‘timeout’ is the total elapsed time for resolution, irrespective of the .attempts and .timeout option values of the configuration.

Returns a `cqueues.dns.packet` answer packet on success, or nil and a numeric error code on failure. The answer may not actually have anything in the ANSWERS section; e.g. if the RCODE is NXDOMAIN.

`resolver:stat()`

Returns a table of statistics for the resolver instance.

field	description
.queries	number of queries submitted
.udp.sent.count	number of UDP packets sent
.udp.sent.bytes	number of UDP bytes sent
.udp.rcvd.count	number of UDP packets received
.udp.rcvd.bytes	number of UDP bytes received
.tcp.sent.count	number of TCP packets sent
.tcp.sent.bytes	number of TCP bytes sent
.tcp.rcvd.count	number of TCP packets received
.tcp.rcvd.bytes	number of TCP bytes received

`resolver:close()`

Explicitly destroy the resolver object, immediately closing all internal descriptors. This routine ensures all descriptors are properly cancelled.

3.2.13 `cqueues.condition`

This module implements a condition variable. A condition variable can be used to queue multiple Lua threads to await a user-defined event. Unlike some condition variable implementations, this one does not implement the monitor pattern directly. A monitor uses both a mutex and a condition variable. However, a full monitor will usually be unnecessary as coroutines do not run in parallel. Monitors are more a necessity in pre-emptive threading environments.

The condition variable primitive can be used to implement mutexes, semaphores, and monitors.

`condition.type(object)`

Returns the string “condition” if the object is a condition variable, or nil otherwise.

`condition.interpose(name, function)`

Add or interpose a condition class method. Returns the previous method, if any.

`condition.new([lifo])`

Returns a new condition variable object. If ‘lifo’ is **true**, waiting threads are woken in LIFO order, otherwise in FIFO order.

Note that the **cqueues** scheduler might schedule execution of multiple woken threads in a different order. The LIFO/FIFO behavior is most useful when implementing a mutex and for whatever reason you wish to select the thread which has waited either the longest or shortest amount of time.

`condition:wait([...])`

Wait on the condition variable. Additional arguments are yielded to the **cqueues** controller for polling. Passing an integer, for example, allows you to effect a timeout. Passing a socket allows you to wait on both the condition variable and the socket.

Returns true if the thread was woken by the condition variable, and false otherwise. Additional values are returned if they polled as ready. It’s possible that both the condition variable and, e.g., a socket object poll ready simultaneously, in which case two values are returned—true and the socket object.

You can also directly yield a condition variable, along with other condition variables, timeouts, or pollable objects, to the **cqueues** controller with **cqueues.poll**.

`condition:signal([n])`

Signal a condition, wakening one or more waiting threads. If specified, a maximum of ‘n’ threads are woken, otherwise all threads are woken.

4 Examples

4.1 HTTP SSL Request

```
1 local cqueues = require"queues"
  local socket = require"queues.socket"
3
  local http = socket.connect("google.com", 443)
5
  local cq = cqueues.new()
7
  cq:wrap(function()
9     http:starttls()

11     http:write("GET_/HTTP/1.0\n")
    http:write("Host:_google.com:443\n\n")
13
    local status = http:read()
15    print("!", status)

17    for ln in http:lines"*h" do
        print("|", ln)
19    end

21    local empty = http:read"*L"
    print "~"
23
    for ln in http:lines"*L" do
25        io.stdout:write(ln)
    end
27
    http:close()
29 end)

31 assert(cq:loop())
```

4.2 Multiplexing Echo Server

```
1 local cqueues = require"queues"
  local socket = require"queues.socket"
3 local bind, port, wait = ...

5 local srv = socket.listen(bind or "127.0.0.1", tonumber(port or 8000))

7 local cq = cqueues.new()

9 cq:wrap(function()
    for con in srv:clients(wait) do
11      cq:wrap(function()
          for ln in con:lines("*L") do
13            cq:write(ln)
              end
15
                cq:shutdown("w")
17      end)
        end
19 end)

21 assert(cq:loop())
```

4.3 Thread Messaging

```
1 local cqueues = require"cqueues"
  local thread = require"cqueues.thread"
3
  -- we start a thread and pass two parameters--'0' and '9'
5 local thr, con = thread.start(function(con, i, j)
    -- the 'cqueues' upvalue defined above is gone
7     local cqueues = require"cqueues"
      local cq = cqueues.new()
9
      cq:wrap(function()
11         for n = tonumber(i), tonumber(j) do
            io.stdout:write("sent_", n, "\n")
13             con:write(n, "\n")
            -- sleep so our stdout writes don't mix
15             cqueues.sleep(0.1)
        end
17     end)

19     assert(cq:loop())
    end, 0, 9)
21

23 local cq = cqueues.new()

25 cq:wrap(function()
    for ln in con:lines() do
27         io.stdout:write(ln, "_rcvd", "\n")
    end
29
    local ok, why = thr:join()
31
    if ok then
33        print(why or "OK")
    else
35        error(require"cqueues.errno".strerror(why))
    end
37 end)

39 assert(cq:loop())
```